

Query Intent Resolver

Engineering closeout | Release decision

No automatic model promotion

The standalone engineering handoff is complete. The production-quality objective is not. Do not automatically promote V1.1: it falsely short-circuited 9 of 62 predicted short circuits (14.52%), versus 2 of 42 (4.76%) for the preserved V1.0.1 reference. Keep the original reference for comparison and V1.1 for research/integration review. Neither is authorized for live traffic.

System / set	Accuracy	False SC / predicted SC	SC recall
V1.0.1 historical reference	78.00%	2/42 (4.76%)	53.33%
V1.1 historical regression	80.33%	9/62 (14.52%)	70.67%
V1.1 core applicability	62.00%	20.00%	64.00%

Completed and reproducible

55 review submissions processed with unchanged originals, separately attributed adjudications and context exclusions. Three fixed model families compared using grouped training-only selection. 120 diagnostic cases executed. 42 unit tests, 134 runtime checks, real HTTP and Docker checks passed. All original V1 files and the 300-query benchmark are preserved.

How to interpret the evidence

The old benchmark has been observed before and is regression evidence, not a new blind test. Training conditions differ. The 100-case core applicability set is assistant-authored; 20 boundary cases are reported separately. The paired accuracy-gain interval includes zero. Neither a successful build nor a favorable point estimate proves production safety.

What remains open

Short-circuit safety improvement remains a model-quality requirement. Peter/MascotGO must confirm actual route handlers, index coverage, invocation surfaces, deployment environment, authorized real-query validation and risk tolerances. No live service or Foundry integration is claimed.

Presentation materials

Repository: Shivansh-Sahni/QueryIntentResolver. Start with v1_1/docs/REVIEWER_BRIEF.html and TECHNICAL_CLOSEOUT.md. Exact evidence is in v1_1/artifacts/SUMMARY.json, RELEASE_DECISION.json and FINAL_VALIDATION.json.